

Databricks Interview Question

How would you optimize a large Delta Lake table?

Consider partitioning, data skipping, Z-Ordering, and file sizes.